





Five-seed model comparison

Frozen validation FPR $\leq 1\%$ thresholds; AE-B uses benign-only calibration

Full dataset | 1,374,143 test flows



Phase 5 matched | 171,753 test targets



Points: mean $\pm$ sample SD (ddof=1), n=5. Test FPR is not constrained to 1%. Cohorts are separate: conventional controls in the right panels are sliced and recalibrated on matched validation targets.